



## The first general-coordinate (QED)(GR)PIC code for astrophysical plasmas

Hayk Hakobyan (Columbia/PPPL)


in collaboration with:

A. Chernoglazov (UMD), B. Crinquand (Princeton → U of Toulouse), A. Galishnikova (Princeton), J. Mahlmann (Columbia),  
A. Philippov (UMD), A. Vanthieghem (Princeton → CNRS), M. Zhou (IAS)

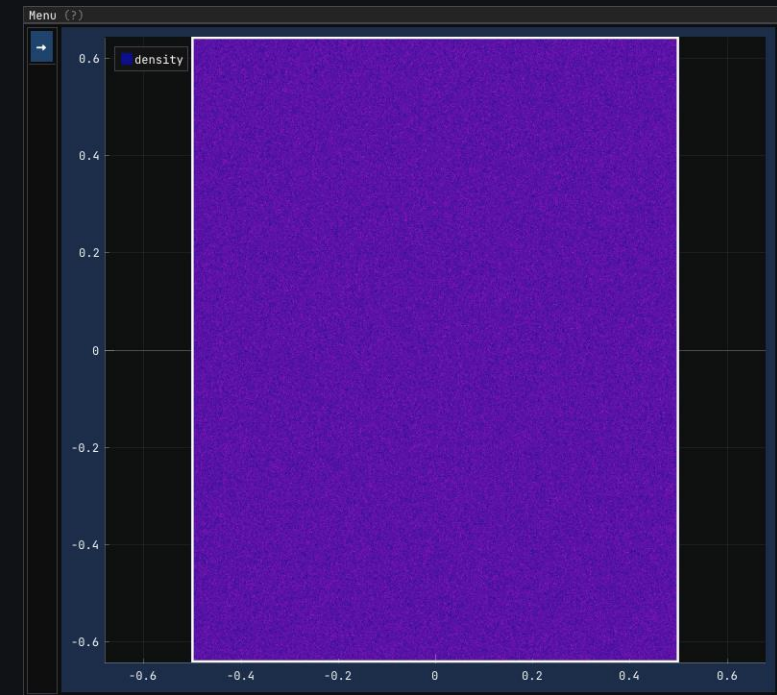
# Capabilities of the *Entity*

@ [haykh.github.io/entity](https://github.com/haykh/entity)

- 1D/2D/3D
- Arbitrary metric/geometry
- General relativity (GR)
- Multi-species
- Radiative processes
- Single-body & two-body QED processes
- Highly modular and flexible

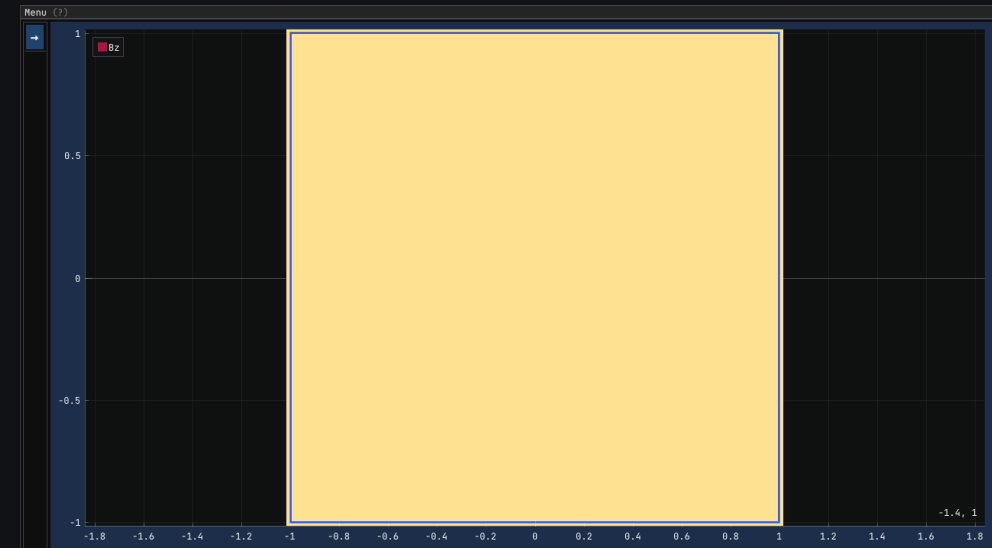


← QED pair-cascade in a neutron star magnetosphere:  
 $e^{\pm}B \rightarrow \gamma$  (curv) then  $\gamma B \rightarrow e^{\pm}$   
(Chernoglazov+, in prep.)



Tearing instability & magnetic reconnection ↑

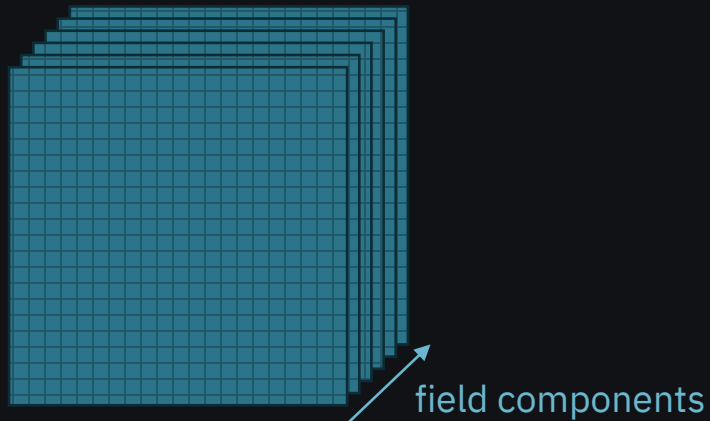
Weibel instability →



# Architecture of the *Entity*


## Memory Layout

```
struct Fields {  
    Kokkos::View<real_t*** [6]> em;  
    Kokkos::View<real_t*** [6]> aux;  
    Kokkos::View<real_t*** [3]> cur;  
    ...  
};
```



```
template <Dimension D>  
struct Metric {  
    KOKKOS_INLINE_FUNCTION  
    void v3_Cntrv2Cov(const coord_t<D>& xi,  
                     const vec_t<D>& v_Cntrv,  
                     vec_t<D>& v_Cov);  
    ...  
};
```

```
struct Particles {  
    float mass;  
    float charge;  
  
    Kokkos::View<int*>    i1;  
    ...  
    Kokkos::View<float*> dx1;  
    ...  
    Kokkos::View<real_t*> ux1;  
    ...  
    Kokkos::View<short*> tag;  
};  
  
auto species = std::vector<Particles> {};
```



# Architecture of the *Entity*

## Kernels

### (a) Field solvers

em(:, :, :, :, :)  
aux(:, :, :, :, :)  
cur(:, :, :, :, :)  
metric

$$\frac{\partial D^i}{c \partial t} = \frac{1}{\sqrt{h}} \varepsilon^{ijk} \partial_j H_k - \frac{4\pi}{c} \frac{J^i}{\sqrt{h}}$$
$$\frac{\partial B^i}{c \partial t} = -\frac{1}{\sqrt{h}} \varepsilon^{ijk} \partial_j E_k$$

$$H_i = \alpha h_{ij} B^j - \frac{1}{\sqrt{h}} \varepsilon_{ijk} \beta^j D^k$$
$$E_i = \alpha h_{ij} D^j + \frac{1}{\sqrt{h}} \varepsilon_{ijk} \beta^j B^k$$

em(:, :, :, :, :)

```
struct FieldSolver_kernel {  
  KOKKOS_INLINE_FUNCTION  
  // vectorized over cells  
  void operator()(int i1, int i2, int i3) const {}  
};
```

### (b) Particle pushers

i1(:, ), ...  
dx1(:, ), ...  
ux1(:, ), ...  
em(:, :, :, :, :)  
metric

$$\frac{du_i}{cdt} = -\gamma \partial_i \alpha + u_j \partial_i \beta^j - \alpha \frac{1}{2\gamma} u_j u_k \partial_i h^{jk} +$$
$$+ \frac{q}{m c^2} \alpha (h_{ij} D^j + \frac{1}{\sqrt{h\gamma}} \varepsilon_{ijk} h^{jl} u_l B^k)$$
$$\frac{dx^i}{cdt} = \frac{\alpha}{\gamma} h^{ij} u_j - \beta^i$$

i1(:, ), ...  
dx1(:, ), ...  
ux1(:, ), ...

```
struct ParticlePusher_kernel {  
  KOKKOS_INLINE_FUNCTION  
  // vectorized over particles  
  void operator()(std::size_t p) const {}  
};
```

# Architecture of the *Entity*

## Kernels

### 3D indexing

```
struct FieldSolver_kernel {
  KOKKOS_INLINE_FUNCTION
  // vectorized over cells
  void operator()(int i1, int i2, int i3) const {
    em(i1, i2, i3, ex1) = ...;
  }
};
```

### 1D indexing

```
struct FieldSolver_kernel {
  KOKKOS_INLINE_FUNCTION
  // vectorized over cells
  // from 0 to nx1 * nx2 * nx3
  void operator()(int cell) const {
    auto [i1, i2, i3] = UNRAVEL(cell, nx1, nx2, nx3);
    em(i1, i2, i3, ex1) = ...;
  }
};
```

[\*] 1D indexing is ~ 20 ... 30% faster  
(on A100 GPU)

# Architecture of the *Entity*

## Kernels

### (c) Current deposition

$$\begin{array}{l} i1(:), \dots \\ dx1(:), \dots \\ i1\_prev(:), \dots \\ dx1\_prev(:), \dots \\ cur(:, :, :, :, :) \end{array} \rightarrow \mathcal{J}^i = \sum_p \frac{q}{\Delta t} (x^i - x_{prev}^i) S(x^i) \rightarrow cur(:, :, :, :, :)$$

```
auto scatter_cur = Kokkos::Experimental::create_scatter_view(cur);
for (auto& sp : species) {
    Kokkos::parallel_for("CurrentsDeposit",
                        sp.npart(),
                        CurrentDeposition_kernel(sp, scatter_cur));
}
Kokkos::Experimental::contribute(cur, scatter_cur);

struct CurrentDeposition_kernel {
    // vectorized over particles
    KOKKOS_INLINE_FUNCTION
    void operator()(std::size_t p) {}
};
```

# Performance

For PIC, the performance is usually evaluated by the **timestep duration per particle**

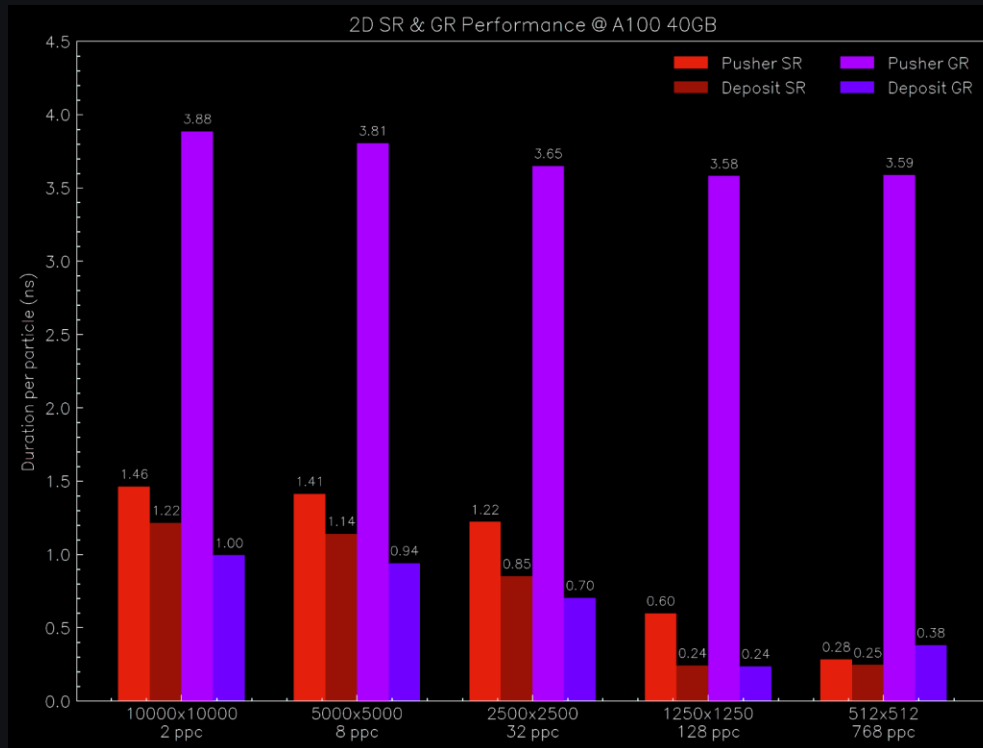
Single core/device performance:

Typical CPU perf.: ~ 20 ... 100 ns/particle (*Tristan v2* @ Intel Cascade Lake)

{ current deposition: ~ 60%  
particle pusher: ~ 30%  
field solver/filtering: ~ 10%

GPU perf.: ~ 0.5 ... 1.5 ns/particle (*Entity* @ NVIDIA A100)


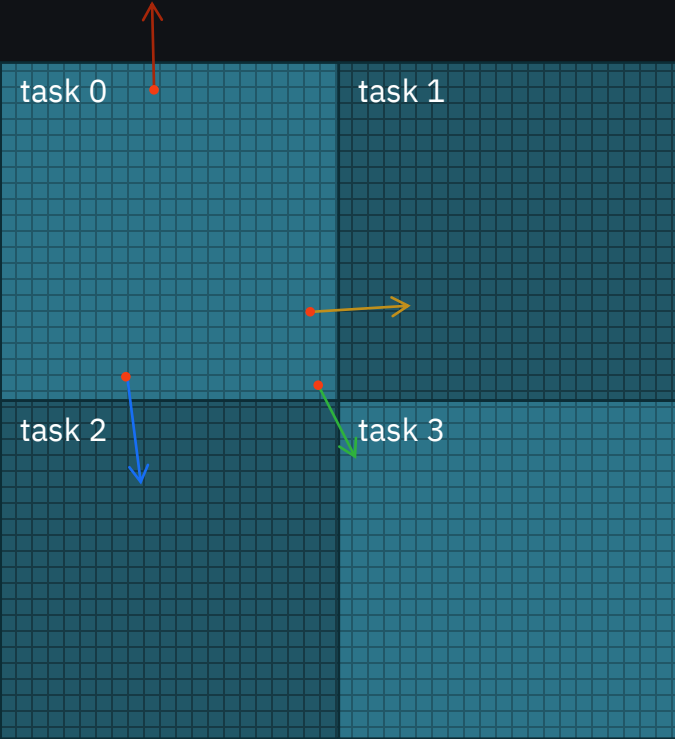
{ current deposition: ~ 35%  
particle pusher: ~ 60%  
field solver/filtering: ~ 5%



~  $2 \cdot 10^8$  particles

# Bottlenecks & Culprits

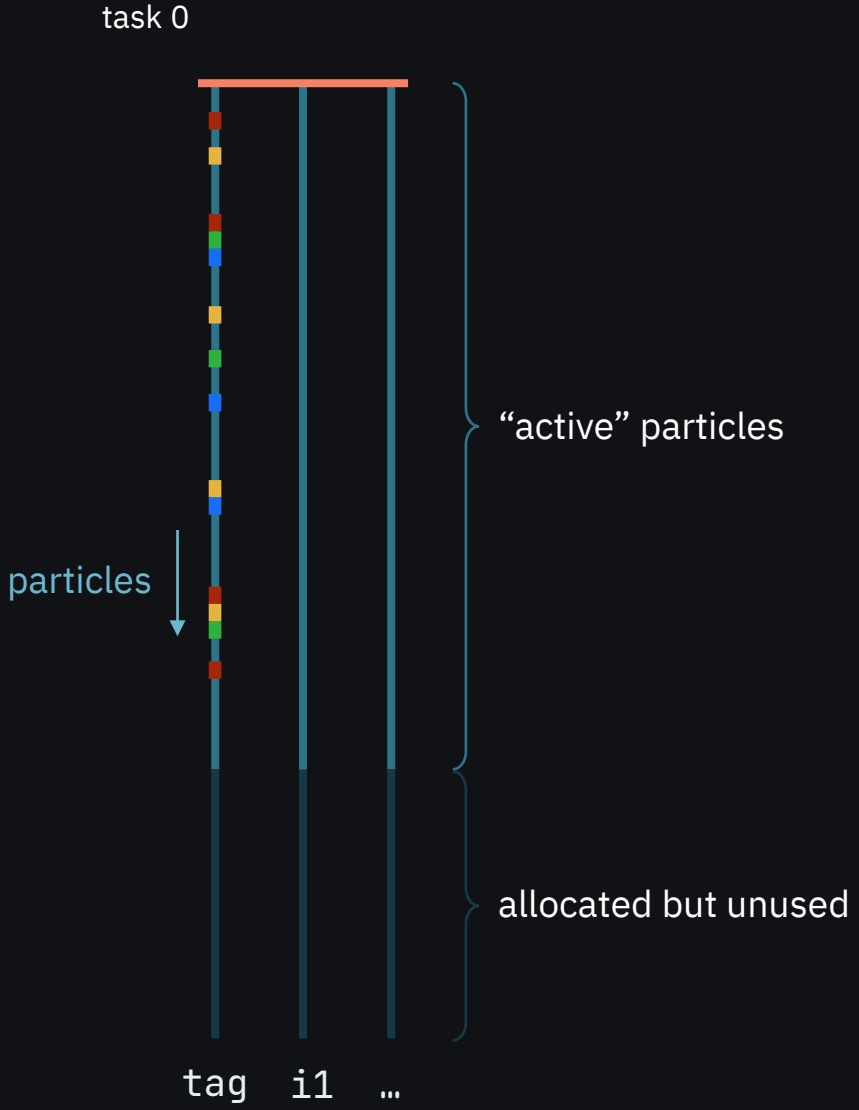
## *MPI communications of particles*





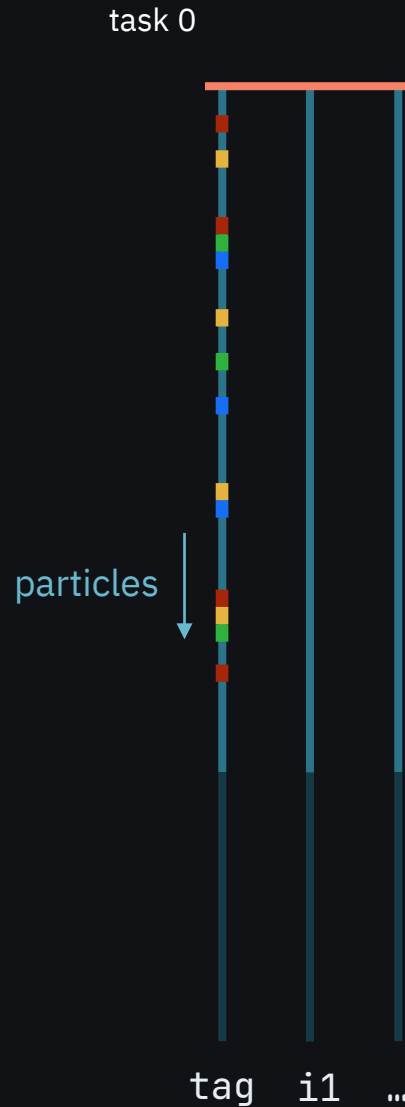
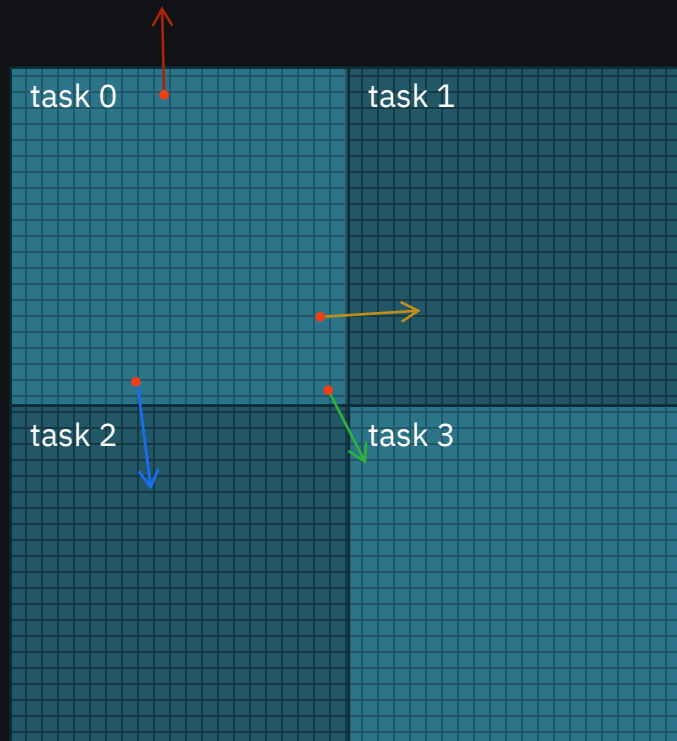
# Bottlenecks & Culprits

## *MPI communications of particles*



# Bottlenecks & Culprits

## MPI communications of particles



```
using KeyType = Kokkos::View<short*>;
using BinOp   = BinTag<KeyType>;
BinOp bin_op(ntags);
Kokkos::BinSort<KeyType, BinOp>
Sorter(Kokkos::subview(prtl.tag, { 0, npart }),
      bin_op,
      false);

Sorter.create_permute_vector();
Sorter.sort(Kokkos::subview(prtl.tag, { 0, npart }));
Sorter.sort(Kokkos::subview(prtl.i1, { 0, npart }));
Sorter.sort(Kokkos::subview(prtl.dx1, { 0, npart }));
Sorter.sort(Kokkos::subview(prtl.ux1, { 0, npart }));
...

template <class KeyViewType>
struct BinTag {
    BinTag(const int& max_bins) : m_max_bins { max_bins } {}

    template <class ViewType>
    inline auto bin(ViewType& keys, const int& i) const -> int {
        return (keys(i) == 0) ? 1 : ((keys(i) == 1) ? 0 : keys(i));
    }


    inline auto max_bins() const -> int {
        return m_max_bins;
    }

    template <class ViewType, typename iT1, typename iT2>
    inline auto operator()(ViewType&, iT1&, iT2&) const -> bool {
        return false;
    }

private:
    const int m_max_bins;
};
```


# Bottlenecks & Culprits

## MPI communications of particles



# Bottlenecks & Culprits

## *MPI communications of particles*



# Bottlenecks & Culprits

## MPI communications of particles

Sorting has to be done every timestep  
⇒ takes up ~ 80% of the timestep ☹

### Possible solutions

1. using *thrust* (similar performance + not portable)

```
thrust::sort_by_key(tag.begin(),  
                  tag.begin() + npart,  
                  thrust::make_zip_iterator(i1, dx1, ux1, ...),  
                  comparator);
```

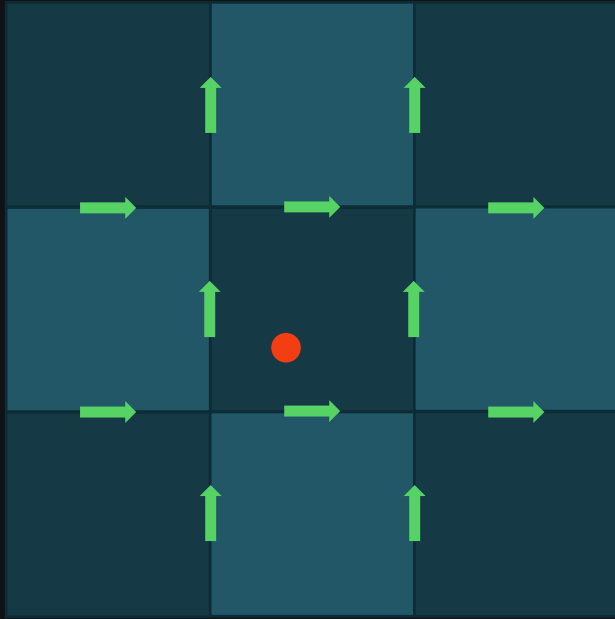
2. pre-extract the data into buffer arrays (not tested)

```
Kokkos::View<std::size_t> send_left_cnt("send_left_cnt");  
"  
Kokkos::parallel_for(  
  "Extract",  
  npart,  
  KOKKOS_LAMBDA(std::size_t p) {  
    if (tag(p) == ParticleTag::sendLeft) {  
      Kokkos::atomic_fetch_add(&send_left_cnt(), 1);  
      send_left_i1(send_left_cnt) = i1(p);  
      tag(p) = ParticleTag::dead;  
    } else if ...  
  });
```

```
using KeyType = Kokkos::View<short*>;  
using BinOp    = BinTag<KeyType>;  
BinOp bin_op(ntags);  
Kokkos::BinSort<KeyType, BinOp>  
Sorter(Kokkos::subview(prtl.tag, { 0, npart }),  
      bin_op,  
      false);  
  
Sorter.create_permute_vector();  
Sorter.sort(Kokkos::subview(prtl.tag, { 0, npart }));  
Sorter.sort(Kokkos::subview(prtl.i1, { 0, npart }));  
Sorter.sort(Kokkos::subview(prtl.dx1, { 0, npart }));  
Sorter.sort(Kokkos::subview(prtl.ux1, { 0, npart }));  
"  
"
```

# Bottlenecks & Culprits

## Field interpolation (gathering)



### Possible solutions

1. sorting by cells
2. reducing the # of intermediate variables
3. storing fields in smaller chunks (tiles)

Example for just the  $E^1$  component (loads 12 values at different locations)

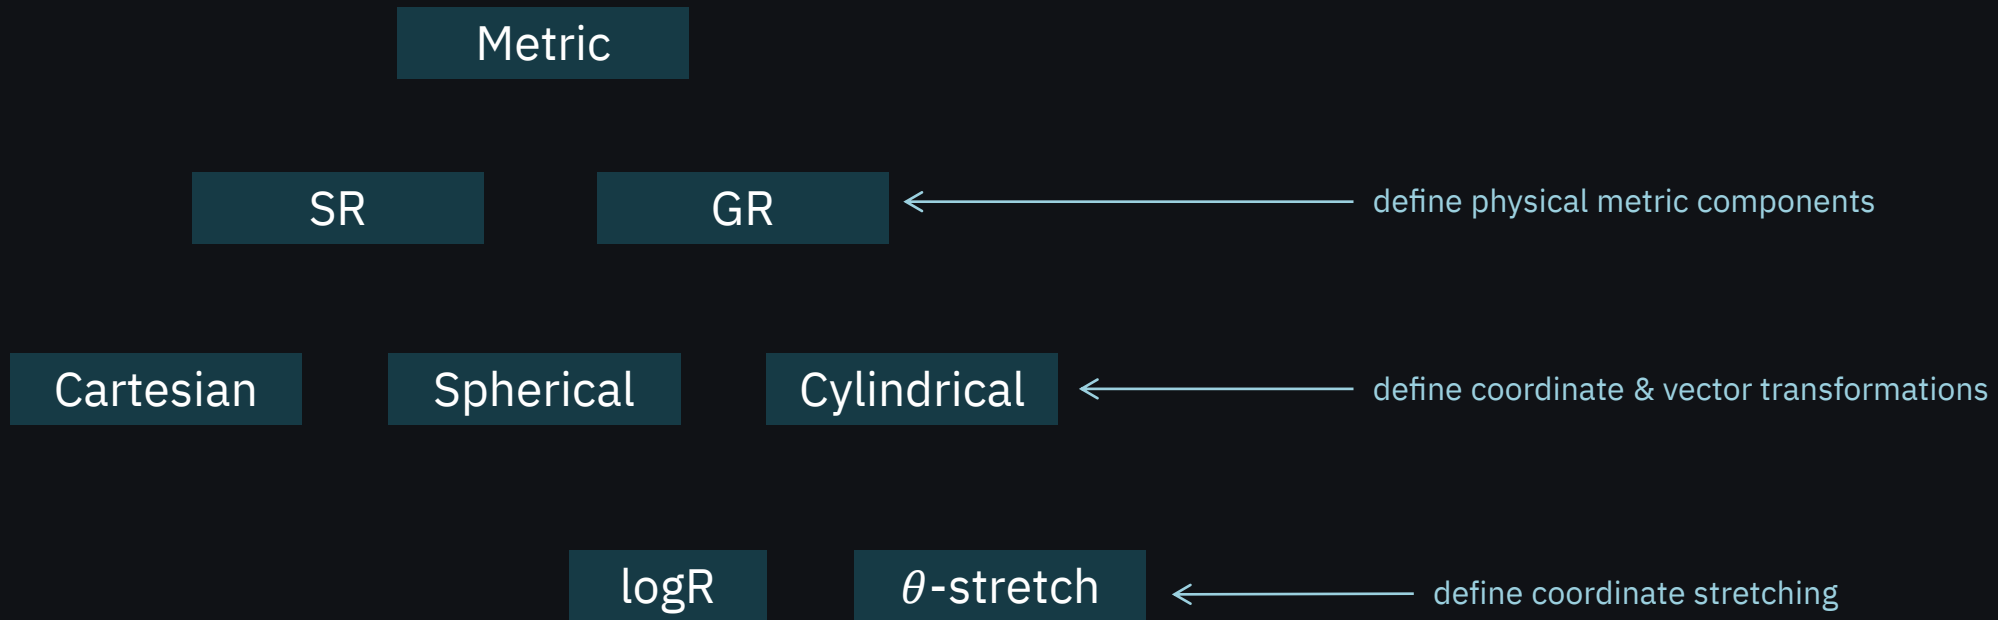
```
struct ParticlePusher_kernel {
    KOKKOS_INLINE_FUNCTION
    // vectorized over particles
    void operator()(std::size_t p) const {
        // interpolate fields ...
    }
};

// Ex1
// interpolate to nodes
c000 = HALF * (EB(i, j, k, em::ex1) + EB(i - 1, j, k, em::ex1));
c100 = HALF * (EB(i, j, k, em::ex1) + EB(i + 1, j, k, em::ex1));
c010 = HALF * (EB(i, j + 1, k, em::ex1) + EB(i - 1, j + 1, k, em::ex1));
c110 = HALF * (EB(i, j + 1, k, em::ex1) + EB(i + 1, j + 1, k, em::ex1));
// interpolate from nodes to the particle position
c00 = c000 * (ONE - dx1_) + c100 * dx1_;
c10 = c010 * (ONE - dx1_) + c110 * dx1_;
c0 = c00 * (ONE - dx2_) + c10 * dx2_;
// interpolate to nodes
c001 = HALF * (EB(i, j, k + 1, em::ex1) + EB(i - 1, j, k + 1, em::ex1));
c101 = HALF * (EB(i, j, k + 1, em::ex1) + EB(i + 1, j, k + 1, em::ex1));
c011 = HALF *
    (EB(i, j + 1, k + 1, em::ex1) + EB(i - 1, j + 1, k + 1, em::ex1));
c111 = HALF *
    (EB(i, j + 1, k + 1, em::ex1) + EB(i + 1, j + 1, k + 1, em::ex1));
// interpolate from nodes to the particle position
c01 = c001 * (ONE - dx1_) + c101 * dx1_;
c11 = c011 * (ONE - dx1_) + c111 * dx1_;
c1 = c01 * (ONE - dx2_) + c11 * dx2_;
ex1_interp = c0 * (ONE - dx3_) + c1 * dx3_;
```

# Bottlenecks & Culprits

## Virtual metric classes

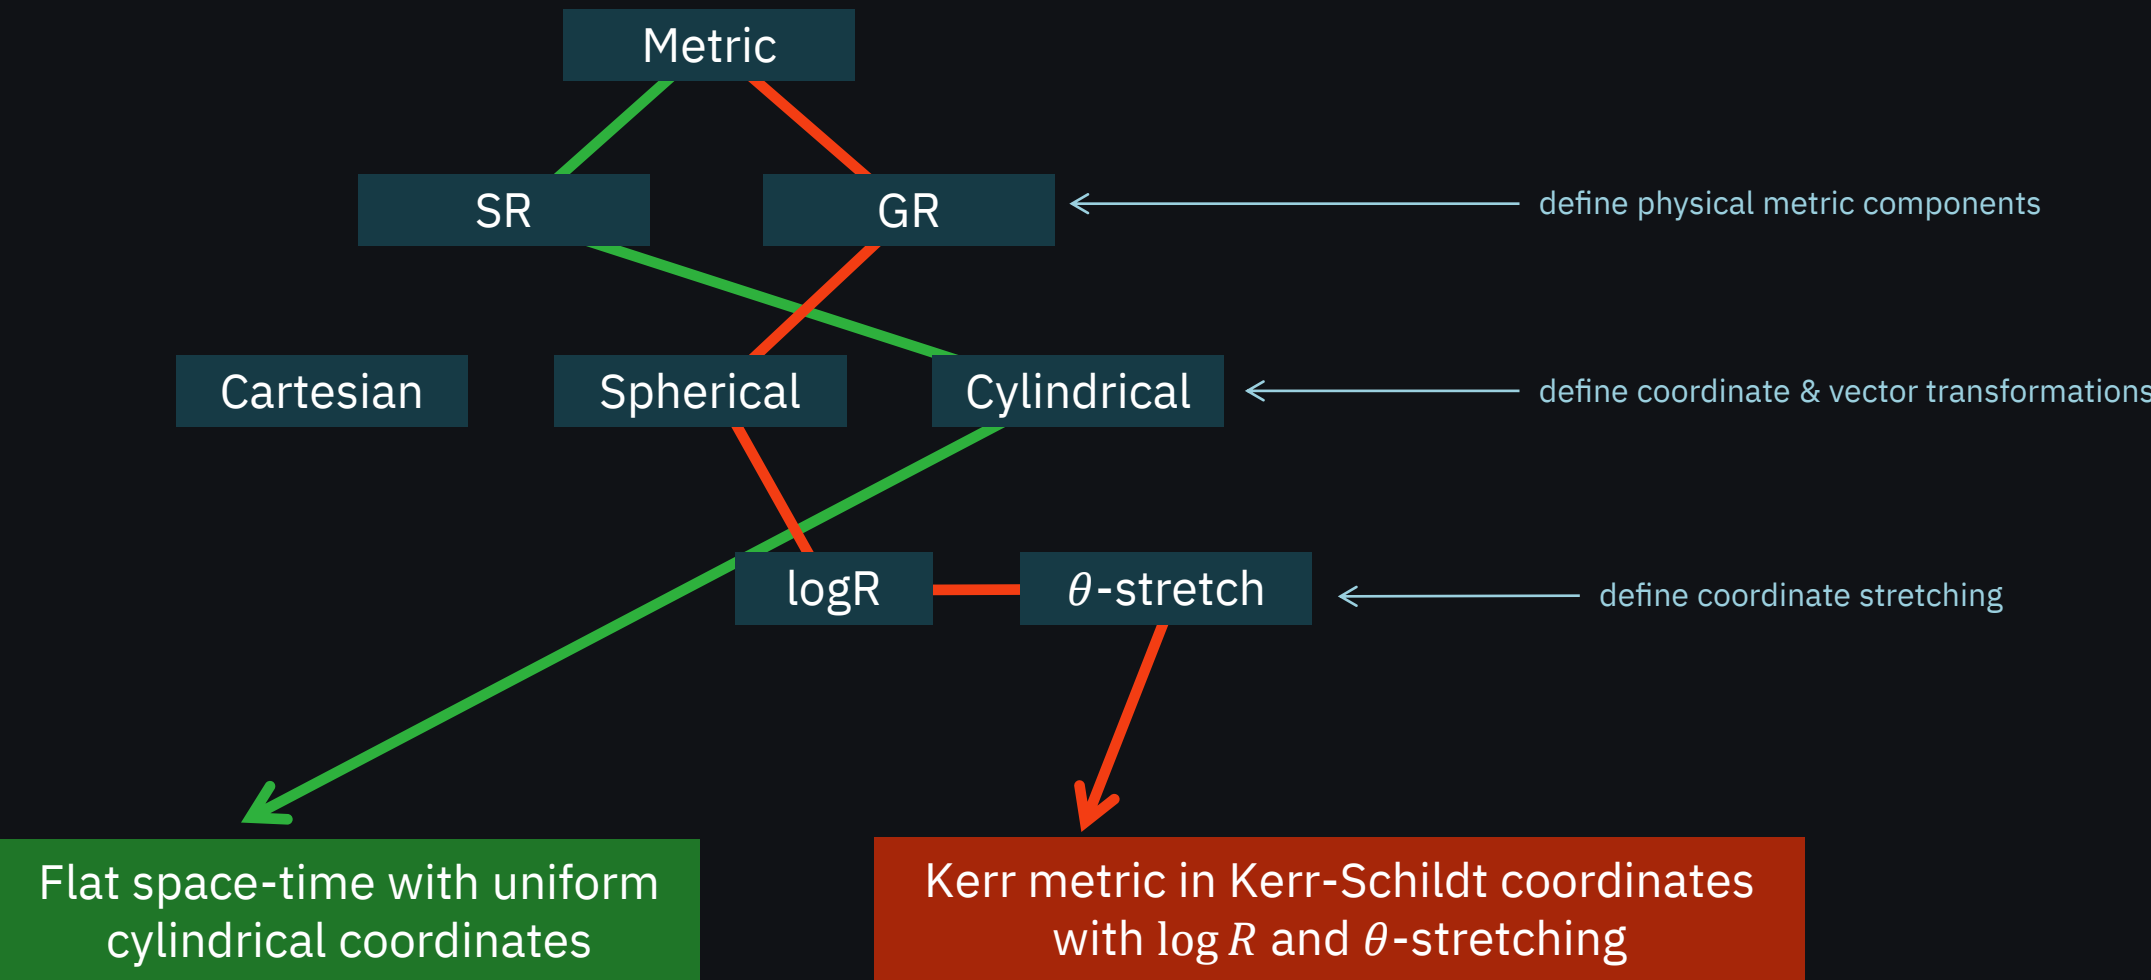
*Entity* is designed to be coordinate (grid) agnostic, so having flexibility in provided metric structures is critical.



# Bottlenecks & Culprits

## Virtual metric classes

Entity is designed to be coordinate (grid) agnostic, so having flexibility in provided metric structures is critical.





# Takeaways

- *Entity* ([haykh.github.io/entity](https://github.com/haykh/entity)) is the first coordinate-agnostic GR PIC code for astrophysical plasmas written with the Kokkos library.
- preliminary performance is fantastic, and we already use the code for actual science!

challenges by priority:

- performance with MPI
- performance of the field gathering
- reformatting the metric structures
- two-body QED module



dev team

- Alexander Chernoglazov [@SChernoglazov](https://github.com/SChernoglazov): PIC}
- Benjamin Crinquant [@bcrinquant](https://github.com/bcrinquant): GR, cubed-sphere}
- Alisa Galishnikova [@alisagk](https://github.com/alisagk): GR}
- Hayk Hakobyan [@haykh](https://github.com/haykh): framework, PIC, GR, cubed-sphere}
- Jens Mahlmann [@jmahlmann](https://github.com/jmahlmann): cubed-sphere, framework, MPI}
- Sasha Philippov [@sashaph](https://github.com/sashaph): all-around}
- Arno Vanthieghem [@vanthieg](https://github.com/vanthieg): PIC, framework}
- Muni Zhou [@munizhou](https://github.com/munizhou): PIC}

We are hiring (ask!)  
<https://jobregister.aas.org/ad/89c8a064>



Supported by the NSF CSSI Grant, and the NVIDIA Corporation Academic Hardware Grant Program.

