

# Kokkos SIMD

Dong Hun Lee, Sandia National Laboratories

Kokkos User Group Meeting 2023

December 14, 2023

Sandia National Laboratories is a multi-mission laboratory managed and operated by National Technology and Engineering Solutions of Sandia, LLC., a wholly owned subsidiary of Honeywell International, Inc., for the U.S. Department of Energy's National Nuclear Security Administration under contract DE-NA0003525.  
SAND2023-14460C

## SIMD

- ▶ Utilization of vector registers in processors to operate on vector lanes in parallel
- ▶ ISA and Intrinsics
- ▶ Compiler auto-vectorization

## Standard C++

- ▶ *C++ Extensions for Parallelism Version 2*<sup>1</sup>
- ▶ *std::simd — merge data -parallel types from the Parallelism TS 2*<sup>2</sup>

---

<sup>1</sup>N4808 Working Draft, C++ Extensions for Parallelism Version 2 by Jared Hoberock

<sup>2</sup>P1928 std::simd - merge data-parallel types from the Parallelism TS 2 by Matthias Kretz

## Kokkos SIMD

- ▶ Abstraction layer of platform-specific vector types for SIMD intrinsics
- ▶ Data-parallel types based on *Extensions for Parallelism, Version 2*<sup>1</sup>
- ▶ Alignment with `std::simd`<sup>2</sup> targeted for ISO standard C++26

---

<sup>1</sup>N4808 Working Draft, C++ Extensions for Parallelism Version 2 by Jared Hoberock

<sup>2</sup>P1928 `std::simd` - merge data-parallel types from the Parallelism TS 2 by Matthias Kretz

	NEON	AVX2	AVX512
float	✓	✓	✓
double	✓	✓	✓
int32_t	✓	✓	✓
uint32_t			✓
int64_t	✓	✓	✓
uint64_t	✓	✓	✓

```
Kokkos::Experimental::simd
```

```
▶ template <class T, class Abi> class simd
```

```
Kokkos::Experimental::simd_mask
```

```
▶ template <class T, class Abi> class simd_mask
```

```
Kokkos::Experimental::native_simd<T>
```

```
Kokkos::Experimental::native_simd_mask<T>
```

native is determined by KOKKOS\_ARCH and Kokkos::DefaultExecutionSpace

## Copy functions

- ▶ `template <class U, class Flags> void copy_from(const U* mem, Flags flags)`
- ▶ `template <class U, class Flags> void copy_to(U* mem, Flags flags)`

## Subscript operators

- ▶ `reference operator [] (std::size_t)`
- ▶ `value_type operator [] (std::size_t)`

- ▶ Basic arithmetic operators
- ▶ Shift operators (logical or arithmetic shifts)
- ▶ Comparison operators
- ▶ Rounding and remainder functions
- ▶ `simd_mask` reductions
- ▶ Subset of `<cmath>` functions

<code>copysign</code>	<code>max</code>	<code>tan</code>	<code>asinh</code>	<code>pow</code>
<code>abs</code>	<code>min</code>	<code>asin</code>	<code>acosh</code>	<code>hypot</code>
<code>sqrt</code>	<code>exp2</code>	<code>acos</code>	<code>atanh</code>	<code>atan2</code>
<code>cbt*</code>	<code>log10</code>	<code>atan</code>	<code>erf</code>	
<code>exp*</code>	<code>log2</code>	<code>sinh</code>	<code>erfc</code>	
<code>log*</code>	<code>sin</code>	<code>cosh</code>	<code>tgamma</code>	
<code>fma*</code>	<code>cos</code>	<code>tanh</code>	<code>lgamma</code>	

- ▶ Not all functions are implemented using intrinsics

*\* requires Intel SVML; otherwise uses a serial fallback implementation*



## Conditionals: where\_expression

- ▶ Provides references to the simd values as described by a mask
- ▶ Used to perform conditional logic using a mask

```
template <class M, class T> class const_where_expression;  
template <class M, class T> class where_expression;  
  
template <class T, class Abi>  
const_where_expression<simd_mask<T, Abi>, simd<T, Abi>>  
where(typename simd<T, Abi>::mask_type const&, simd<T, Abi> const&);  
  
template <class T, class Abi>  
where_expression<simd_mask<T, Abi>, simd<T, Abi>>  
where(typename simd<T, Abi>::mask_type const&, simd<T, Abi>&);
```

**Conditionals:** `where_expression`

## Copy functions

- ▶ `copy_from`
- ▶ `copy_to`
- ▶ `gather_from`
- ▶ `scatter_to`

## Assignment operators

- ▶ Assignment
- ▶ Compound assignments

```
double x[N];  
double y[N];  
double z[N];  
double r[N];  
  
for (int i = 0; i < N; ++i) {  
    r[i] = sqrt(x[i] * x[i] + y[i] * y[i] + z[i] * z[i]);  
}
```

```
#include <Kokkos_SIMD.hpp>

using simd_type = Kokkos::Experimental::native_simd<double>;
using tag_type = Kokkos::Experimental::element_aligned_tag;

constexpr int width = int(simd_type::size());
simd_type simd_x, simd_y, simd_z, simd_r;

for (int i = 0; i < N; i += width) {
    simd_x.copy_from(x + i, tag_type());
    simd_y.copy_from(y + i, tag_type());
    simd_z.copy_from(z + i, tag_type());
    simd_r = Kokkos::sqrt(simd_x * simd_x + simd_y * simd_y + simd_z * simd_z);
    simd_r.copy_to(r + i, tag_type());
}
```

```
#include <Kokkos_SIMD.hpp>

using simd_type      = Kokkos::Experimental::native_simd<double>;
using simd_view_type = Kokkos::View<simd_type*>;
using tag_type       = Kokkos::Experimental::element_aligned_tag;

constexpr int simd_view_length = N / simd_type::size();

simd_view_type simd_vx("simd_view_x", simd_view_length);
/*...*/

Kokkos::RangePolicy<> policy(0, simd_view_length);
Sqrt_Functor sqrt_functor(simd_vx, simd_vy, simd_vz, simd_vr);
Kokkos::parallel_for(policy, [&](int i) {
    int offset = i * simd_type::size();
    simdv_x(i).copy_from(x + offset, tag_type());
    simdv_y(i).copy_from(y + offset, tag_type());
    simdv_z(i).copy_from(z + offset, tag_type());
    sqrt_functor(i);
    simdv_r(i).copy_to(r + offset, tag_type());
});
```

## Currently unavailable operations

- ▶ `simd` compound assignments
- ▶ `simd` reductions
- ▶ `simd` constructor taking a mem pointer as an argument