

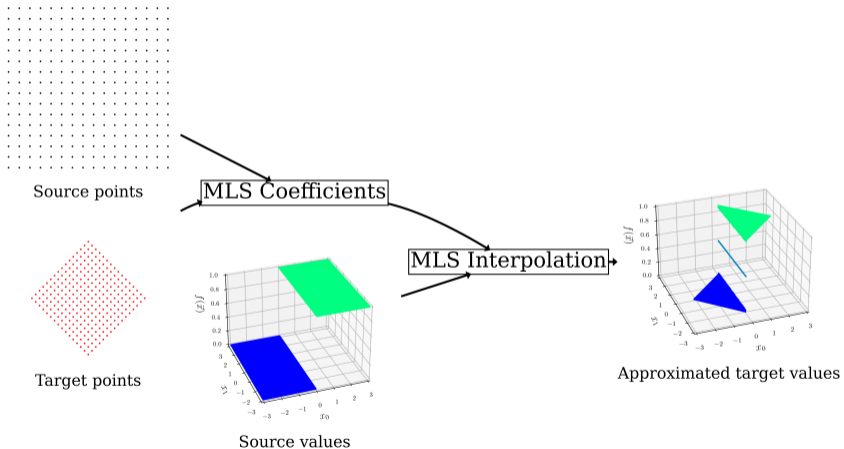
KUG 2023 - Moving least squares

Yohann Bosqued
December 14, 2023

ORNL is managed by UT-Battelle LLC for the US Department of Energy

Moving least squares

Principle



Definitions

- Let $A(\vec{x})$ be a polynomial, then $A(\vec{x}) = p(\vec{x}) \cdot a^T = \langle p(\vec{x}) | a \rangle$ with $p(\vec{x})$, the polynomial basis taken at \vec{x} .

$$p\left(\begin{bmatrix} x & y \end{bmatrix}\right) = \begin{bmatrix} 1 & x & y & x^2 & xy & y^2 \end{bmatrix}$$

- Let ϕ be a compactly supported radial basis function. Its value is only dependant on the input's norm and is supported on a compact set. Its goal is to select points and give each an influence on the final result.

$$\phi(\vec{x}) = \begin{cases} (1 - \|\vec{x}\|)^2 & \text{if } \|\vec{x}\| \leq 1 \\ 0 & \text{else} \end{cases}$$

Weighted and moving least squares

Let f be a function from \mathbb{R}^d to \mathbb{R} and a set of source points S . The goal of the weighted least squares is to find the polynomial G minimizing $\|G - f\|_{\phi,2}$. Using the polynomial basis, finding g minimizing $\|\langle p(\cdot)|g \rangle - f\|_{\phi,2}$.

$$\begin{aligned}g^T &= [P^T \Phi P]^{-1} P^T \Phi F^T \\f(\vec{u}) &\approx \langle p(\vec{u})|g \rangle \\&\approx p(\vec{u}) [P^T \Phi P]^{-1} P^T \Phi F^T\end{aligned}$$

Weighted and moving least squares

Let f be a function from \mathbb{R}^d to \mathbb{R} and a set of source points S . The goal of the weighted least squares is to find the polynomial G minimizing $\|G - f\|_{\phi,2}$. Using the polynomial basis, finding g minimizing $\|\langle p(\cdot)|g \rangle - f\|_{\phi,2}$.

$$\begin{aligned}g^T &= [P^T \Phi P]^{-1} P^T \Phi F^T \\f(\vec{u}) &\approx \langle p(\vec{u})|g \rangle \\&\approx p(\vec{u}) [P^T \Phi P]^{-1} P^T \Phi F^T\end{aligned}$$

$$\begin{aligned}g_{\vec{u}}^T &= [P_{\vec{u}}^T \Phi_{\vec{u}} P_{\vec{u}}]^{-1} P_{\vec{u}}^T \Phi_{\vec{u}} F_{\vec{u}}^T \\f(\vec{u}) &\approx \langle p(\vec{u})|g_{\vec{u}} \rangle \\&\approx p(\vec{u}) [P_{\vec{u}}^T \Phi_{\vec{u}} P_{\vec{u}}]^{-1} P_{\vec{u}}^T \Phi_{\vec{u}} F_{\vec{u}}^T\end{aligned}$$



$$\begin{aligned}f(\vec{u}) &\approx \langle G_{\vec{u}}|F_{\vec{u}} \rangle \\G_{\vec{u}} &= p(\vec{u}) [P_{\vec{u}}^T \Phi_{\vec{u}} P_{\vec{u}}]^{-1} P_{\vec{u}}^T \Phi_{\vec{u}}\end{aligned}$$

What needs to be done



- From a set of points S and a target point \vec{u} :
 - Find the closest set of neighboring points $S_{\vec{u}}$
 - Compute and save the coefficients ($G_{\vec{u}} = \rho(\vec{u}) [P_{\vec{u}}^T \Phi_{\vec{u}} P_{\vec{u}}]^{-1} P_{\vec{u}}^T \Phi_{\vec{u}}$)
 - Now ready to make approximations!
- Coefficients and approximation are computed in batches of target points



Example (Coefficients)





Example (Coefficients)




- 6 neighbors per target point

Example (Coefficients)



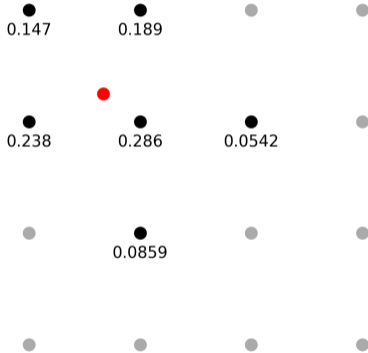

- Linear poly basis (size 3)
- $P_{\vec{u}}: 6 \times 3$
- $\Phi_{\vec{u}}: 6 \times 6$

Example (Coefficients)



- $M_{\vec{u}}: 3 \times 3$
- $M_{\vec{u}}^{-1}: 3 \times 3$


Example (Coefficients)



- $G_{\vec{u}}: 6 \times 1$

Example (Approximation)


$$f(\vec{x}) = \frac{x_0 x_1}{4} + 1$$




$$\begin{bmatrix} 0.147 \\ 0.189 \\ 0.238 \\ 0.286 \\ 0.054 \\ 0.086 \end{bmatrix} \cdot \begin{bmatrix} 1 \\ 1.75 \\ 1 \\ 1.5 \\ 2.0 \\ 1.25 \end{bmatrix}$$

Example (Larger scale)


$$f(\vec{x}) = \frac{\text{sgn}(x_0)+1}{2}$$



Source values



Real target values



Approximated values



Interface



```
Kokkos::View<Point*, MemorySpace> source_points;  
Kokkos::View<Point*, MemorySpace> target_points;  
Kokkos::View<double*, MemorySpace> source_values;  
Kokkos::View<double*, MemorySpace> approx_values;  
  
// ...  
  
MovingLeastSquares<MemorySpace> mls(exec_space, source_points, target_points);  
mls.interpolate(exec_space, source_values, approx_values);  
// source_values = ...  
mls.interpolate(exec_space, source_values, approx_values);
```

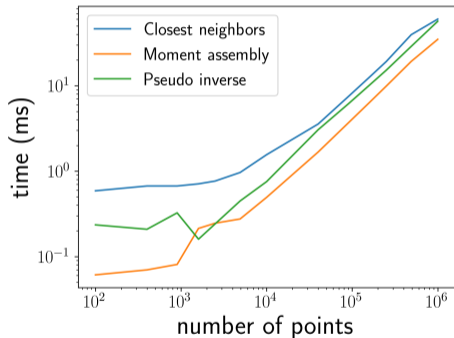
Performance

Performance (1M points)

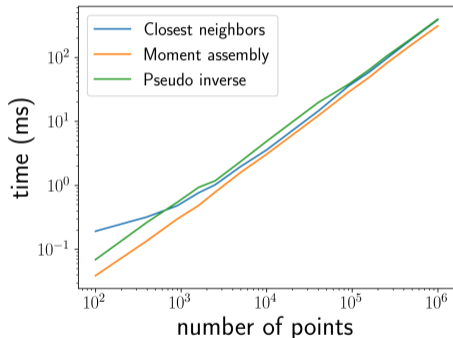


Performance



Cuda



Serial





Performance (Update!)



- No scratch pad use, global memory only
- Multiple kernels, heavy use of MDRangePolicy

Performance (Update!)



- Scratch pad for temporary data
- Single kernel, parallelized with TeamPolicy (One target per thread)
- Originally 1200B/target, now down to 960B/target with extra work

Learning Kokkos

Cursus and previous internships



- Parallelization
- OpenMP
- MPI
- ...

- CUDA
- OpenMP Target
- GPU offloading
- ...

What I used to do



```
#pragma omp target distribute teams parallel reduce op(+:...) is_device_ptr(...)  
for (int i = 0; i < N; i++) {  
    // ...  
}
```

and manual memory management

What I do now

```
Kokkos::parallel_reduce(Kokkos::RangePolicy<ExecutionSpace>(space, 0, N),  
KOKKOS_LAMBDA (int i, double& loc) {  
    // ...  
}, /**/);
```

and smart memory management (with Views!)

What will I do



Questions?