

Kokkos: State on Exascale Architectures

Daniel Arndt, Oak Ridge National Laboratory

Kokkos User Group Meeting 2023

December 12, 2023

This manuscript has been authored by UT-Battelle, LLC, under contract DE-AC05-00OR22725 with the U.S. Department of Energy. The United States Government retains and the publisher, by accepting the article for publication, acknowledges that the United States Government retains a nonexclusive, paid-up, irrevocable, world-wide license to publish or reproduce the published form of this manuscript, or allow others to do so, for United States Government purposes.

Primary Programming Models: OpenMP, Cuda, HIP, SYCL



LANL Crossroads
Intel CPUs,
OpenMP, SYCL?



LBNL Perlmutter
AMD CPU, NVIDIA GPU
CUDA



ORNL Frontier
AMD CPU, AMD GPU
HIP



ANL Aurora
Intel CPUs, Intel GPUs
SYCL



LLNL EI Capitan
AMD CPU, AMD GPU
HIP

Supercomputer

- ▶ Summit(ORNL)
- ▶ Perlmutter(LBNL)
- ▶ Sierra(LLNL)
- ▶ Polaris(ANL)

Minimum version Cuda 11.0.0.

```
cmake ..\  
-DCMAKE_CXX_COMPILER=* \  
-DKokkos_ENABLE_CUDA=ON \  
-DKokkos_ARCH_NATIVE=ON \  
-DKokkos_ARCH_AMPERE80=ON
```

- ▶ Without compiler wrapper, clang++ or nvcc_wrapper must be used.
- ▶ nvhpc only used as host compiler by default.

Results from `bytes_and_flops(TeamPolicy)`¹

scalar	Bandwidth	Compute	Cache
float	1251 GiB/s	14280 GFlop/s	3762 GiB/s
double	1267 GiB/s	7592 GFlop/s	6938 GiB/s
int32_t	1222 GiB/s	18457 GFlop/s	4684 GiB/s
int64_t	1267 GiB/s	3778 GFlop/s	6895 GiB/s

- ▶ Peak FP64 Vector: 19.5 TFLOPS
- ▶ Memory Bandwidth: 1.6 TB/sec
- ▶ Cache Size: L1/L2: 192KB (per SM)/40 MB

¹This research used resources of the Argonne Leadership Computing Facility, a U.S. Department of Energy (DOE) Office of Science user facility at Argonne National Laboratory and is based on research supported by the U.S. DOE Office of Science-Advanced Scientific Computing Research Program, under Contract No. DE-AC02-06CH11357.

Supercomputer

- ▶ Frontier(ORNL)
- ▶ El Capitan(LLNL)

Unsupported features

- ▶ Tasks

Minimum version ROCm 5.2.0.

```
cmake ..\  
  -DCMAKE_CXX_COMPILER=hipcc \  
  -DKokkos_ENABLE_HIP=ON \  
  -DKokkos_ARCH_NATIVE=ON \  
  -DKokkos_ARCH_AMD_GFX90A=ON
```

Results from bytes_and_flops (TeamPolicy)²

scalar	Bandwidth	Compute	Cache
float	1160 GiB/s	20544 GFlop/s	2756 GiB/s
double	1140 GiB/s	19320 GFlop/s	2883 GiB/s
int32_t	1150 GiB/s	20194 GFlop/s	2757 GiB/s
int64_t	1140 GiB/s	4979 GFlop/s	2865 GiB/s

- ▶ Peak FP64 Vector: 23.95 TFLOPS
- ▶ Memory Bandwidth: 1.6 TB/sec
- ▶ Cache Size: L1/L2: 16KB (per CU)/16 MB

²This research used resources of the Oak Ridge Leadership Computing Facility at the Oak Ridge National Laboratory, which is supported by the Office of Science of the U.S. Department of Energy under Contract No. DE-AC05-00OR22725.

When using `HIPManagedSpace`, the memory migrates between the CPU and the GPU if:

- ▶ the hardware supports it and
- ▶ the kernel was compiled to support page migration and
- ▶ the environment variable `HSA_XNACK` is set to 1

L1 cache much smaller than on other GPUs. Decreasing memory usage via Enabling `Kokkos_ENABLE_HIP_MULTIPLE_KERNEL_INSTANTIATIONS` might improve performance but increases compilation time.

Using `-munsafe-fp-atomics` forces hardware-based floating-point atomics (no synchronization across kernels).

Supercomputer

- ▶ Aurora(ANL)

Unsupported features

- ▶ WorkGraphPolicy
- ▶ Tasks
- ▶ Graphs (dummy implementation)
- ▶ Virtual functions/function pointer
<https://github.com/intel/llvm/pull/10540>
- ▶ `::printf` → `Kokkos::printf`


```
cmake ..\  
-DCMAKE_CXX_COMPILER=icpx \  
-DKokkos_ENABLE_SYCL=ON \  
-DKokkos_ARCH_NATIVE=ON \  
-DKokkos_ARCH_INTEL_PVC=ON
```

- ▶ Replace last line with `-DKokkos_ARCH_INTEL_GEN=ON` for JIT compilation.
- ▶ Minimum version oneAPI 2023.0.0.

Results from bytes_and_flops(TeamPolicy)³

scalar	Bandwidth	Compute	Cache
float	1002 GiB/s	17484 GFlop/s	4973 GiB/s
double	960 GiB/s	8746 GFlop/s	6928 GiB/s
int32_t	1007 GiB/s	6108 GFlop/s	4714 GiB/s
int64_t	958 GiB/s	982 GFlop/s	4715 GiB/s

- ▶ Peak FP64 Vector: 22.9 TFLOPS/tile
- ▶ Memory Bandwidth: 1.6 TB/sec/tile
- ▶ Cache Size: 128KB (per work group)/408 MB

³This research used resources of the Argonne Leadership Computing Facility, a U.S. Department of Energy (DOE) Office of Science user facility at Argonne National Laboratory and is based on research supported by the U.S. DOE Office of Science-Advanced Scientific Computing Research Program, under Contract No. DE-AC02-06CH11357.

For `Kokkos::RangePolicy` with `Kokkos::parallel_for`, the workgroup size can be manually specified:

```
Kokkos::parallel_for(  
  Kokkos::RangePolicy<ExecutionSpace>(space, 0, N)  
    .set_chunk_size(1024),  
  *this);
```

The subgroup size can be forced via

```
export IGC_ForceOCLSIMDWidth=32
```

or as compiler flag

```
-fsycl-default-sub-group-size=32
```

Also, see <https://github.com/kokkos/kokkos/pull/6496>, that would allow

```
Kokkos::RangePolicy<ExecSpace, Kokkos::SubGroupSize<16>>
```

Larger amount of registers can be requested via

```
export SYCL_PROGRAM_COMPILE_OPTIONS=\
    "-ze-opt-large-register-file"
```

or as compiler flag

```
-Xs "-options -ze-opt-large-register-file"
```

useful if there are spills

[...] compiled SIMD32 allocated 128 regs and spilled around 4

Using <https://github.com/kokkos/kokkos/pull/5879> allows specifying the scratch level at compile-time and helps with

[...] warning: Adding 4 occurrences of additional control flow due to presence of generic address space operations

Multidimensional View access produces poor assembly. Replace

```
for (int i=0; i<N; ++i)
    view(i,3,10) = ...;
```

with

```
auto stride = view.stride(0);
auto view_ptr = &view(0,3,10);
for (int i=0; i<N; ++i, view_ptr+=stride)
    *view_ptr = ...;
```

Questions?