

# Kokkos Usage in xRAGE

*Peter Maginot- EAP DPL for Physics*  
with Daniel Holladay, Zach Medin, Clell (CJ) Solomon

Kokkos User Group Meeting, December 12-15, 2023

LA-UR-23-33701

# Outline

1. What is xRAGE
2. Porting history/strategy
3. FY24 plans
4. Kokkos usage within grey diffusion

# xRAGE is a large LANL ASC multiphysics code

- Mostly Fortran 90/95
- O(500k) lines
- Large number of internal users
  - Their day job is to run xRAGE
- O(15-20) developers
- Extensive validation basis
- MPI domain decomposition for parallelism
- Expected to run [performantly] on all large NNSA HPC systems
- Geared toward high energy density applications (NIF, Omega, Z experiments)
  - But still has “cold” physics for HE burn, material strength, etc..

**Interested in Kokkos for GPU porting because code base is too large to maintain multiple code paths**

# xRAGE targeting of GPUs has not been steady

- Significant work in FY19 and FY20
  - Focus on inter-operatbility, FLCL, DualView
  - Start porting physics [folders]
- Impediments
  - COVID
  - Unilateral decision to stop porting
  - LANL ASC re-organization, EAP re-organization
  - Staff departures (Classification Office, Industry, Other projects → Industry)
- Resumption in late FY22 through Today
  - Focus on verifying / demonstrating GPU performance of ported packages

**We are not where we could be, but we have a consistent vision and plan to move forward!**

# Each ported xRAGE physics currently responsible for migrating data to and from the GPU

- The above also implies, “from Fortran to C++”. Currently:
  1. Break apart Fortran derived types into component arrays and scalar
  2. Convert arrays into `flcl_ndarray_t` objects
  3. Cross the Fortran/C barrier
  4. Convert `flcl_ndarray_t*` objects into Host Views
  5. Transfer Host Views to Device Views
- Very manual, very large function signatures
- FLCL has caused some tricky to find issues
  - `FLCL::HostSpace` not really `HostSpace` (it's `CudaUVM`)
  - Correctness of taking a Fortran allocated array and assuming it is in `CudaUVM` space?
  - `flcl_ndarray_t` have served their purpose, but new options exist

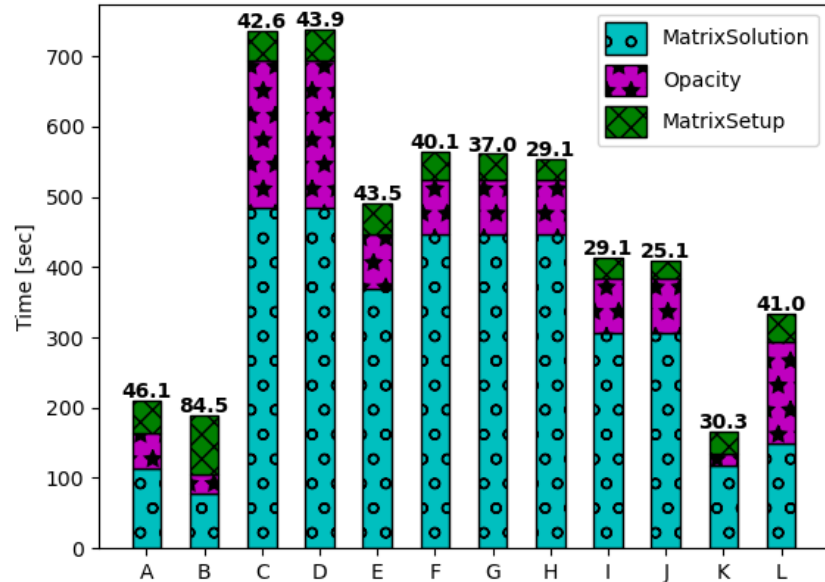
# FY23 was “Year of the Machine” FY24 will be xRAGE’s “Year of the GPUs”

- Unsplit hydrodynamics, conduction, high explosives reactive burn, multigroup radiation diffusion to be ported
  - LANL is getting serious since impending Venado and El Cap hardware deliveries suggest GPUs are not going to go away
- Removing FLCL dependency
  - xRAGE mixed compiler build + availability of “newer” GNU compilers on Sierra
  - xRAGE will make use of F2018 “C Descriptors” to replace `flcl_ndarray_t`
- Streamline steps from Fortran derived types to GPU Views
  - Maintain historical hierarchy of derived type data
  - Consistent method to auto-generate Fortran to C interoperable structs then transform
    - Most developers will write a single Python file per Fortran derived type
    - CMake will run Python files to 1) auto-generate interoperable structs definitions and 2) C++ transforms to get from C descriptors to Views

# [Grey] radiation diffusion is our most ported physics

- Amongst the original three “physics” packages [folders] identified in 2018
- Three pieces for complete grey diffusion solve
  - Opacity (data) lookup
  - Matrix setup
  - Matrix solve
- Matrix setup is effectively the piece of xRAGE physics we can control
  - Results suggest we might wish to consider controlling more pieces
- Since it was most ported, radiation diffusion has been our Kokkos testing ground
  - What features of Kokkos might we want to use for performance?
  - Allows us to explore how do multiple GPU chefs in the kitchen work?
    - xRAGE Kokkos, CUDA Fortran, CUDA C, TPLs that use Kokkos all must build and link together in harmony

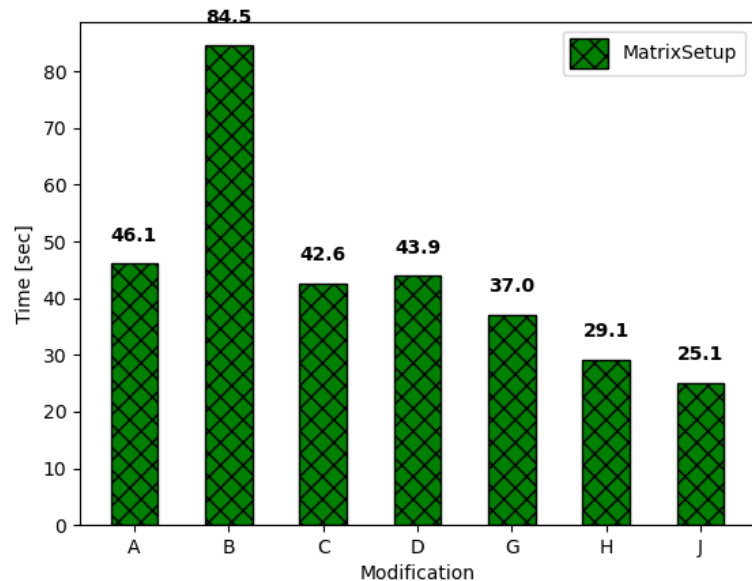
# Despite following later, xRAGE is experiencing much of others' progression in performance



	CPU	Improvement
A	36-CTS	Legacy Fortran
B	40-P9	C++ (forAllFaces)
C	4-P9	no MPS
D	4-P9	forAllFaceType
E	4-P9	lag opacity
F	4-P9	hyre 2.26 (bad settings)
G	4-P9	partition_space streams
H	4-P9	cudaMallocAsync
I	4-P9	hyre 2.26 (better settings)
J	4-P9	unmanaged memory
K	36-CTS	Final C++ (lag opacity)
L	36-CTS	Final C++ (no opacity lagging)



# Focusing on what we are actually speeding up paints a fairer picture



	CPU	Improvement
A	36-CTS	Legacy Fortran
B	40-P9	GPU + MPS
C	4-P9	no MPS
D	4-P9	forAllFaceType splitting
G	4-P9	partition_space streams
H	4-P9	cudaMallocAsync
J	4-P9	unmanaged memory

We can achieve modest speed-up of a low FLOP, high data transfer routine via code re-writing and less common Kokkos strategies.

# xRAGE iteration over faces needed to evolve for GPUs

- Faces in xRAGE are categorized by type and facing direction
  1. Lo boundary
  2. Hi boundary
  3. Interior at 1:1 interface
  4. Interior fine [lo] / coarse [hi]
  5. Interior coarse [hi] / fine [lo]
- Face-based data stored in 2-D arrays
  - Global face index not a concept
  - Indexed within a dimension
  - <10% over allocation ( $\text{maxint} * \text{numdim}$ )

$\text{maxint} = \max(\text{max\_f\_x}, \text{max\_f\_y}, \text{max\_f\_z})$

## Fortran Iteration Pattern

```
do dim=1,numdim
  do loop=1, n_types_of_faces(dim)
    face_type = face_id(loop,dim)
    if face_type == 1
      n_lo = face_lo(loop,dim)
      n_hi = face_hi(loop,dim)
      do n=n_lo,ni
        cell_hi = face_local(n,HI,dim)
        face_data(n,dim) = func(data(cell_hi))
      enddo
    else if face_type == 2
      ! ... omitted for brevity
    enddo
  enddo
enddo
```

**Natively would require 15 kernel launches!**

# Radiation diffusion first unrolled common pattern once and created auxiliary iteration structures

- Did not collect timing data prior to transition
  - Code did not work on GPUs prior to a MR that did too many things 😊
- Pros
  - Single kernel launch
  - Logical “forAllFaces” of a Fortran pattern
- Cons:
  - Increased memory footprint
  - Retains face\_type checking logic
  - Atomics needed for reductions to cell data

## forAllFaces Iteration Pattern

```
parallel_for("calc_a_thing",
  RangePolicy<EXEC_SPACE>(0, f_idata.n_faces_tot),
  KOKKOS_LAMBDA (const size_t f_idx) {
    cell_lo = f_data.cell_lo_of_face(f_idx);
    cell_hi = f_data.cell_hi_of_face(f_idx);
    face_type = f_data.type_of_face(f_idx);
    dim = f_data.dim_of_face(f_idx);
    idx = f_data.idx_of_face(f_idx);
    if (face_type == 2){
      face_data(n,dim) = func(data(cell_lo))
    }
    // ... omit other face_types for brevity
  });
```

# forAllFaces has been split into 3 separate kernel launches

- Eliminated if checking
  - FaceliterationData doubles in memory footprint
- Small latency slowdown without `Kokkos::Experimental::partition_space`
- Streams allowed for
  - simultaneous kernels
  - less `cudaDeviceSynchronize` calls
- Streams increased complexity
  - As code exits a function, leave streams in flight
  - Manual process for book keeping

## forAllFaceTypes Iteration Pattern

```
parallel_for("calc_a_thing_interior",
  RangePolicy<EXEC_SPACE>
  (streams[0],0,f_idata.n_faces_int),
  KOKKOS_LAMBDA (const size_t f_idx) {
    // ...
  });
parallel_for("calc_a_thing_typ1",
  RangePolicy<EXEC_SPACE>
  (streams[1],0,f_idata.n_faces_typ1),
  KOKKOS_LAMBDA (const size_t f_idx) {
    // ...
  });
parallel_for("calc_a_thing_typ2",
  RangePolicy<EXEC_SPACE>
  (streams[2],0,f_idata.n_faces_typ2),
  KOKKOS_LAMBDA (const size_t f_idx) {
    // ...
  });
```

xRAGE operator splitting severely limits the amount of overlapping computation that can occur

# Streaming of View allocation + enabling cudaMallocAsync resulted in the single largest speed-up of run\_diff\_cycle

- xRAGE is used to “free” allocations with CPU+DDR
  - More than 120 allocations within run\_diff\_cycle sized proportional to spatial DOF
- Streams permitted early computations to overlap with View creation on device
  - Further complicated code flow / readability
    - Juggling 11(!) streams
- Change in behavior from 3.7.01 to 4.0.01
  - Our data lookup functions are all on the CPU  
`create_mirror_view_and_copy(view_alloc(stream[0], HOST(), WithoutInitializing, dev_view)`
- Requests
  - Please no more static assert failures that don't give a line number
  - Spack variant in Kokkos maintained package.py for cudaMallocAsync

# cudaMallocAsync still noticeable when calling 100+ times

- Hand rolled a memory pool for exploratory purposes
  - 120+ allocations to 6
- Umpire will replace manual pointer math
  - Can Kokkos handle Umpire allocators being used to evict / transfer data?
  - Can underlying pointer of Unmanaged views be swapped?

```
n_big_alloc += local_count + mirror_count + copy_in_count;
View<double*> big_alloc(view_alloc(STORE(), "big_alloc" , stream[1])
,
                        n_big_alloc);
stream[1].fence();
double* big_alloc_ptr = big_alloc.data();
size_t big_alloc_used = 0;
double* tev_ptr = big_alloc_ptr + big_alloc_used;
View<double*,MemoryUnmanaged> tev (tev_ptr, hv_tev.size());
big_alloc_used += tev.size();
deep_copy(stream[1],tev, hv_tev);
```

# xRAGE's path to performance is porting more physics

- Grey diffusion matrix setup has been a testbed for Kokkos concepts to improve performance
  - `Kokkos::Experimental::partition_space`
  - `cudaMallocAsync`
  - `Kokkos::MemoryUnmanaged`
- Would like to see `partition_space` come out of `Experimental`
- Considering moving to 4.1 for profiling concurrent kernels
- Interested in seeing if others have
  - Explored/use memory pools and/or whether it is a priority for Kokkos development
  - Have Power9 + V100 results comparing ScatterView vs. atomics for reductions

# Questions / Comments / Advice?