

Kokkos Kernels: Overview

Luc Berger-Vergiat, Sandia National Laboratories

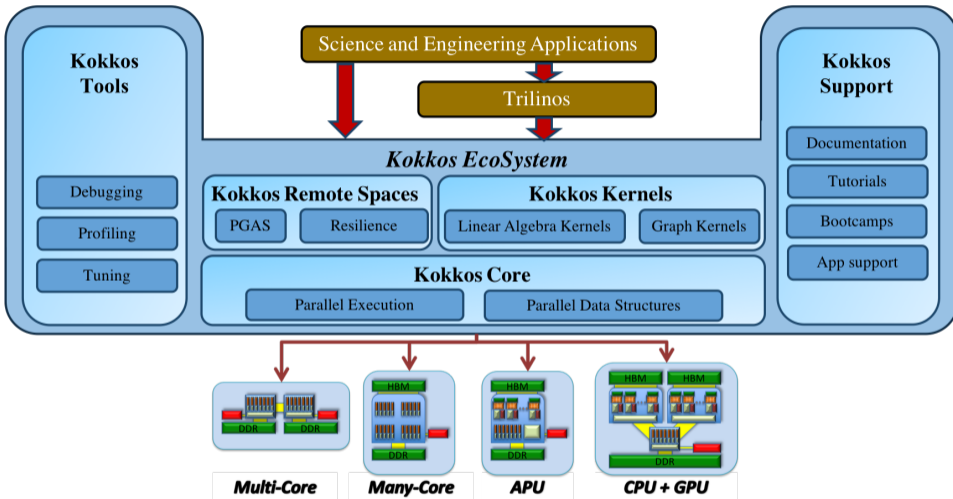
Kokkos User Group Meeting 2023

December 12, 2023

Sandia National Laboratories is a multi-mission laboratory managed and operated by National Technology and Engineering Solutions of Sandia, LLC., a wholly owned subsidiary of Honeywell International, Inc., for the U.S. Department of Energy's National Nuclear Security Administration under contract DE-NA0003525.
SAND2023-14345C

Overview

Kokkos Kernels in the Kokkos Ecosystem



Kokkos Kernels team:

- ▶ Siva Rajamanickam *team lead*
- ▶ Luc Berger-Vergiat *team co-lead*
- ▶ Vinh Dang
- ▶ Nathan Ellingwood
- ▶ James Foucar
- ▶ Brian Kelley
- ▶ Kim Liegeois
- ▶ Carl Pearson
- ▶ Ernesto Prudencio

- ▶ Batched (dense and sparse)
- ▶ BLAS
- ▶ Graph
- ▶ **Lapack** → New host base interface primarily wrapping TPLs (LAPACK, cuSOLVER, rocSOLVER, MAGMA)
- ▶ **ODE** → Device based implementation of Runge-Kutta and BDF time integrators
- ▶ Sparse

Latest development

BLAS completeness (no band/packed implementation)

- ▶ BLAS 1: complete
- ▶ BLAS 2: complete
- ▶ BLAS 3: need SYMM, HEMM and rank k/2k updates

Execution Space interface (stream/queue execution) supported for all kernels

- ▶ `KokkosBlas::myBlasKernels(const ExecutionSpace& space, ...)`

SYCL backend support

- ▶ Expending oneAPI MKL support, building/testing on Aurora

Example calling BLAS `gemm` with an execution space instance to run the kernel in a stream.



```
using execution_space = Kokkos::Cuda;

auto instances =
    Kokkos::Experimental::partition_space(execution_space(), 1, 1);
KokkosBlas::gemm(instances[0], tA, tB, alpha, A1, B1, beta, C1);
KokkosBlas::gemm(instances[1], tA, tB, alpha, A2, B2, beta, C2);
Kokkos::fence(); // All results available after this point.
```


- ▶ Sparse format conversion: `coo2csr`, `csc2csr`
- ▶ SpGEMM supports reuse and has improved TPL support (MKL, cuSPARSE, rocSPARSE)
- ▶ New incomplete factorization algorithms
 - ▶ `parILUt` algorithm (iterative computation of L and U)
 - ▶ `MDF(0)` reorders matrix rows to minimize discarded fill
 - ▶ Stream version of `ILU(k)` and `SpTRSV`
- ▶ `CrsMatrix` sort and merge (also graph version)
- ▶ `SpMV` improved for `BsrMatrix`

Sparse Batched algorithms

- ▶ Algorithms implemented
 - ▶ Linear algebra (SpMV, Vector operations...)
 - ▶ Iterative solver (CG, GMRES)
 - ▶ Preconditioner (Jacobi)
- ▶ Launch parameters tuned by architecture
 - ▶ NVIDIA V100
 - ▶ AMD MI50/MI250



New component for time integration algorithms (still experimental)

- ▶ Explicit integrators
 - ▶ Runge Kutta (orders 1 to 5)
 - ▶ Various schemes for stability (Fehlberg 45, Cash-Karp, Dormand-Prince)
 - ▶ Time adaptive
 - ▶ DOP 853 upcoming...
- ▶ Implicit integrators
 - ▶ BDF (orders 1 to 5)
 - ▶ Time and order adaptive implementation
 - ▶ Adams-Moulton method upcoming...
- ▶ Methods are implemented to be called in innermost parallel level (SerialInternal)

Upcoming work

Short term goals:

- ▶ Block version of ILU(K)
- ▶ Adding BLIS TPL for BLAS implementation
- ▶ LAPACK features and TPL support will expand greatly
- ▶ Improve stream preconditioners: balancing, reordering...
- ▶ Expend oneAPI MKL utilization

Longer term goals:

- ▶ automated architecture based tuning
- ▶ more performance monitoring using benchmarks
- ▶ documentation improvements
- ▶ handle refactor
- ▶ full BLAS layer

Currently Kokkos Kernels is primarily developed at Sandia.

More external collaborations are welcomed

- ▶ Join us on kokkosteam.slack.com in the `#kokkos-kernels` channel
- ▶ Create an issue or pull request on github.com/kokkos/kokkos-kernels
- ▶ If desirable, an external collaboration meeting can be setup

Conclusion

Thank you for your attention! Any questions?

