

# Kokkos and MLIR

Brian Kelley and Kim Liegeois, Sandia National Laboratories

Kokkos User Group Meeting 2023

December 12, 2023

Sandia National Laboratories is a multi-mission laboratory managed and operated by National Technology and Engineering Solutions of Sandia, LLC., a wholly owned subsidiary of Honeywell International, Inc., for the U.S. Department of Energy's National Nuclear Security Administration under contract DE-NA0003525.  
SAND2023-14543C

# Introduction

- ▶ MLIR
- ▶ Why MLIR is useful to the Kokkos ecosystem
- ▶ Previous Work
- ▶ Examples: ResNet18 and SpMV
- ▶ Ongoing Work
- ▶ Conclusion

# Overview



- ▶ MLIR: Multi-Level Intermediate Representation
- ▶ Part of the LLVM project
- ▶ Like LLVM IR, an SSA language designed to be automatically analyzed and optimized by a compiler
- ▶ Unlike LLVM, includes high-level operations, e.g. matrix multiplication
- ▶ Operations organized into families called **dialects**

- ▶ Some built-in dialects, higher to lower level:
  - ▶ `linalg`: high-level tensor, matrix and vector operations
  - ▶ `scf`: structured control flow: parallel loops, `if`, `while`
  - ▶ `memref`: allocate and access multidimensional arrays
  - ▶ `gpu`: heterogeneous memory and GPU kernels in a CUDA-like model
  - ▶ `llvm`: LLVM IR instructions (making MLIR a superset of LLVM)
- ▶ MLIR includes **passes** which analyze and transform a program
  - ▶ Lowering: convert high-level operation(s) to lower-level equivalent
  - ▶ Optimizing: e.g. loop fusion, tiling, strength reduction...
- ▶ Users can also create their own dialects and passes

# Motivation



## Why bring Kokkos and MLIR together?

- ▶ Interfacing between Python ML frameworks (PyTorch etc.) and Kokkos C++ HPC codes
  - ▶ Direct interfacing: increased developer effort, and introduced dependencies on Python packages
  - ▶ Instead, MLIR can be used to automatically generate C++ source code from Python ML models. Integrate this into Kokkos-based HPC codes.



Why bring Kokkos and MLIR together?

- ▶ Automation of tedious programming tasks
  - ▶ Host-device memory migration
  - ▶ Parameter search and autotuning
  - ▶ Overlapping of independent computations using execution space instances

## Previous Work

- ▶ torch-mlir: external open-source project that compiles PyTorch models to MLIR
  - ▶ Includes end-to-end infrastructure to compile, load and run a model within Python
- ▶ Kokkos → MLIR emitter: 2022 exploratory project at Sandia
  - ▶ Project investigated whether MLIR was a suitable interface between Kokkos and machine learning frameworks
  - ▶ Emitter generates Kokkos C++ source code from MLIR program with mid-level dialects
  - ▶ Successfully emitted ResNet18 pre-trained model, then compiled and ran with CUDA backend
- ▶ Sparse tensor dialect
  - ▶ Effort led by Dr. Aart Bik (Google) to add first-class sparse tensor support to MLIR
  - ▶ Describe tensor formats and layouts at high level: CRS, block structure, doubly-compressed dimensions, etc.
  - ▶ Parallel code generation from high-level `linalg` operations

# Examples

- ▶ ResNet18 (CNN image classifier) end-to-end example.
- ▶ Begin with pre-trained PyTorch model:

```
ResNet(  
  (conv1): Conv2d(3, 64, ...)  
  (bn1): BatchNorm2d(64, ...)  
  (relu): ReLU(inplace=True)  
  (maxpool): MaxPool2d(kernel_size=3, ...)  
  (layer1): Sequential(...)  
  ...
```

- ▶ Use torch-mlir to generate high-level MLIR
- ▶ All weights arrays are included in the MLIR as constant `memrefs`
- ▶ Then use built-in passes to lower to mid-level dialects
- ▶ These dialects most closely match Kokkos's level of abstraction

- ▶ %N: an SSA value (the result of an operation)
- ▶ scf.parallel: parallel N-dimensional loop
- ▶ memref.load: read a value from an array
- ▶ arith.divf: floating point division

```
scf.parallel (%arg1, %arg2, %arg3, ...) {  
  %249 = memref.load %224[%arg1, %arg2]  
  %250 = arith.divf %249, %cst : f32  
  memref.store %250, %225[%arg1, ...]  
  scf.yield  
}
```

- ▶ Use the Kokkos emitter to generate C++ code, one operation at a time

```
Kokkos::parallel_for(Kokkos::MDRangePolicy<
exec_space, Kokkos::Rank<4>>(…),
KOKKOS_LAMBDA(int64_t unit_v1255, int64_t …)
{
    int64_t v1255 = v10 + unit_v1255 * v7;
    int64_t v1256 = v10 + unit_v1256 * v7;
    int64_t v1257 = v10 + unit_v1257 * v7;
    int64_t v1258 = v10 + unit_v1258 * v7;
    float v1259 = v249(v1255, v1256);
    float v1260 = v1259 / v2;
    …
}
```



- ▶ Compile the C++ code to a shared library, with Kokkos linked in
- ▶ Load the library into Python with CTypes, and run model inference

```
# Example.py
import kokkosModule
predictions = kokkosModule.forward(image)
```

image:



```
predictions:
[('Labrador retriever', 70.657),
 ('golden retriever', 4.988), ...]
```

- ▶ Or, use the C++ code directly from an existing Kokkos program
- ▶ Takes image as input (RGB,  $224 \times 224$ ) and returns probability vector (1000 classes)

```
// mlir_kokkos_module.cpp
Kokkos::View<float[1][1000], Kokkos::LayoutRight>
forward(Kokkos::View<float[1][3][224][224], Kokkos::LayoutRight> v1) {
    ...
}
```

- ▶ Use PyTACO (included with MLIR) to express tensor formats and computation
- ▶ Use sparse tensor dialect to generate parallel sparse matrix times vector kernel

```
A = pt.tensor([5, 5], [pt.dense, pt.compressed], dtype=pt.float64)
b = pt.tensor([A.shape[1]], [pt.dense], dtype=pt.float64)
c = pt.tensor([A.shape[0]], [pt.dense], dtype=pt.float64)
# fill A and b...
i, j = pt.get_index_vars(2)
c[i] = A[i,j] * b[j]
```

- ▶ Apply sparse tensor lowering pipeline
- ▶ Use Kokkos emitter to generate kernel with two-level parallelism

```
void kokkos_sparse_kernel_0(
  Kokkos::View<double*, ...> v1, // c values
  Kokkos::View<size_t*, ...> v2, // A row offsets
  Kokkos::View<size_t*, ...> v3, // A column entries
  Kokkos::View<double*, ...> v4, // A values
  Kokkos::View<double*, ...> v5 // b values
) {
  typedef Kokkos::TeamPolicy<exec_space>::member_type member_type;
  int league_size = (5 - 0 + 1 - 1) / 1;
  Kokkos::TeamPolicy<exec_space> policy (league_size, Kokkos::AUTO(), Kokkos::AUTO() );
  Kokkos::parallel_for(policy, KOKKOS_LAMBDA(member_type member)
  {
    int64_t unit_v6 = member.league_rank ();
    int64_t v6 = 0 + unit_v6 * 1;
    ...
    Kokkos::parallel_reduce(
      Kokkos::TeamVectorRange(member, (v10 - v8 + 1 - 1) / 1),
      [&](const int64_t &unit_v12, double& v13)
      {
        ...
        double v17 = v15 * v16;
        v13 += v17;
      }, v11)
    // memref.store
    v1(v6) = v11;
    ...
  }
}
```

## Ongoing and Future Work

- ▶ Tensor partitioning and distributed computations
  - ▶ We have a prototype `partition` dialect
  - ▶ Can express a 2D block distributed SpMV now
- ▶ Use partitioning infrastructure to target spatial dataflow accelerators
  - ▶ These systems have many small processors in a grid, with fast communication between neighbors
  - ▶ Want to distribute sparse tensors and computations over processors
  - ▶ Have successfully run SpMV kernel on NextSilicon Maverick using OpenMP programming model
- ▶ Automatic differentiation at the MLIR level

# Conclusion

- ▶ MLIR can express ML and HPC programs at varying levels of abstraction
- ▶ Transformations can exploit high-level information that would not be available to a conventional C++ compiler
- ▶ We are developing tools with MLIR to benefit HPC use cases:
  - ▶ MLIR → Kokkos C++ emitter: convert MLIR to portable Kokkos-based source code
  - ▶ Partition dialect: support tiled and distributed sparse tensors
  - ▶ Target novel spatial dataflow accelerators
- ▶ **Thank you! Any questions?**