

BENCHMARKING PORTABLE STAGGERED FERMION KERNEL WRITTEN IN KOKKOS AND MPI¹

KOKKOS USER GROUP MEETING 2023

12th December, 2023 | Simon Schlepphorst | Jülich Supercomputing Centre

¹DOI: [10.1145/3624062.3624179](https://doi.org/10.1145/3624062.3624179)

Staggered fermions

staggered fermionic action

$$S_F = a^4 \sum_{n \in \Lambda} \bar{X}(n) \left(\sum_{\mu=1}^4 \eta_{\mu}(n) \frac{U_{\mu}(n)X(n + \hat{\mu}) - U_{\mu}^{\dagger}(n - \hat{\mu})X(n - \hat{\mu})}{2a} + mX(n) \right)$$

- space-time coordinates $n = (n_1, n_2, n_3, n_4)$ separated by a lattice spacing a
- and $U_{\mu} \in \mathbf{C}^{3 \times 3}$, $X \in \mathbf{C}^3$
- with direction index μ
- required dimensions for naive layout $[n_1][n_2][n_3][n_4][4][3][3][2]$
- using $L = n_1 = n_2 = n_3$, $T = n_4$ and $L = T$ for convenience

Staggered fermions

staggered fermionic action

$$S_F = a^4 \sum_{n \in \Lambda} \bar{X}(n) \left(\sum_{\mu=1}^4 \eta_{\mu}(n) \frac{U_{\mu}(n)X(n + \hat{\mu}) - U_{\mu}^{\dagger}(n - \hat{\mu})X(n - \hat{\mu})}{2a} + mX(n) \right)$$

arithmetic intensity (for 32-bit floating point numbers)

$$I = \frac{570 \text{ FLOP}}{792 \text{ B}} = 0.72 \text{ FLOP/B.}$$



Implementation in Kokkos and MPI

Data Layout

```
using complex_t = Kokkos::complex<float>;  
using Vector_t = Kokkos::View<complex_t ****[3]>;  
using Link_t = Kokkos::View<complex_t ****[4][3][3]>;
```

- use in place buffer for computation

$$L^3 \cdot T = 32^4$$

```
Vector_t Y(34,34,34,34);  
Vector_t X(34,34,34,34);  
Link_t U(33,33,33,33);
```

- and additional continuous buffer for communication with MPI
- apply local periodic boundary conditions by recalculating the site indices

Implementation in Kokkos and MPI

- different execution spaces for bulk and halo

Execution Spaces

```
using BulkSpace_t = Kokkos::DefaultExecutionSpace;
```

```
using HaloSpace_t = Kokkos::DefaultExecutionSpace;
```

```
BulkSpace_t BulkExecSpace = BulkSpace_t();
```

```
HaloSpace_t HaloExecSpace = HaloSpace_t();
```

- simple MDRange

Execution Policy

```
auto PB = Kokkos::MDRangePolicy<Kokkos::Rank<4>>(BulkSpace, {1,1,1,1}, {33,33,33,33});
```

```
auto PH = Kokkos::MDRangePolicy<Kokkos::Rank<4>>(HaloSpace, {1,1,1,1}, {2,33,33,33});
```

Implementation in Kokkos and MPI

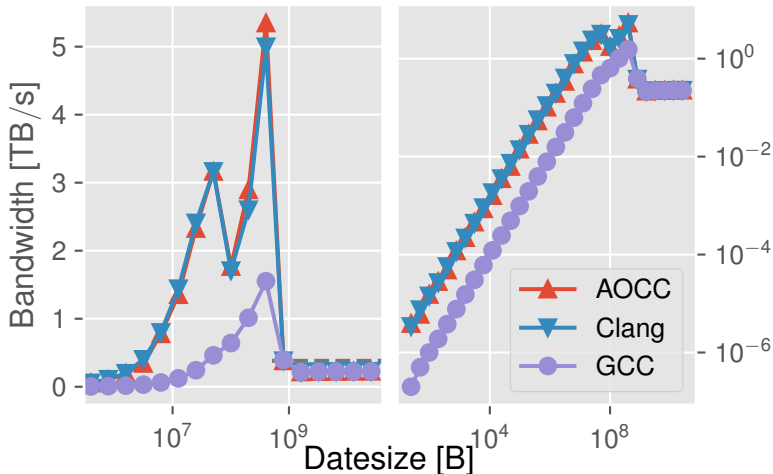
Kernel Algorithm (Input: U , X Output: Y)

```
 $n \in \Lambda$   
for  $i \leftarrow 1, 2, 3$  do  
   $t \leftarrow 0$   
  for  $j \leftarrow 1, 2, 3$  do  
    for  $\mu \leftarrow 1, 2, 3, 4$  do  
       $t \leftarrow t + U_\mu(n)_{ij} * X(p(n + \hat{\mu}))_j$   
       $t \leftarrow t - U_\mu^*(p(n - \hat{\mu}))_{ji} * X(p(n - \hat{\mu}))_j$   
    end for  
  end for  
   $Y_i \leftarrow t$   
end for
```

- $p()$ calculates the correct n according to periodic boundaries

AMD Ryzen 7742 (x86 CPU, Dual Socket)


Stream Triad



JURECA @ JSC, Kokkos 3.6, AOCC 3.2, Clang 13.0, GCC 11.2

AMD Ryzen 7742 (x86 CPU, Dual Socket)


Staggered



JURECA @ JSC, Kokkos 3.6, AOCC 3.2, Clang 13.0, GCC 11.2

Fujitsu A64FX (ARM CPU)

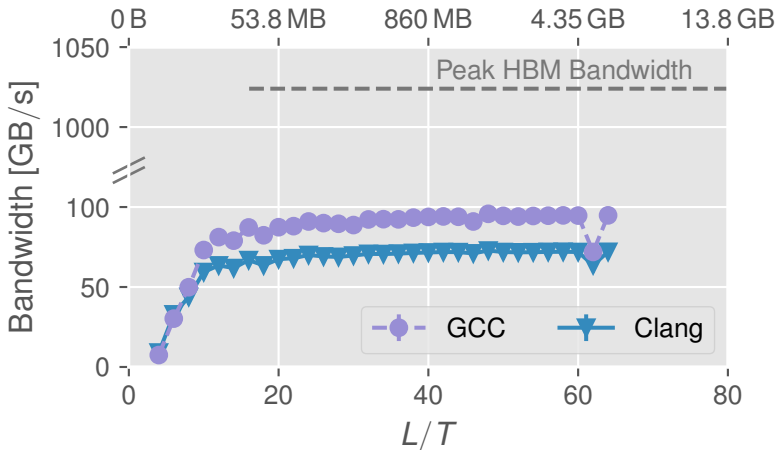
Stream Triad



CTE-ARM @ BSC, Kokkos 3.6, GCC 11.1, Clang 14.0

Fujitsu A64FX (ARM CPU)

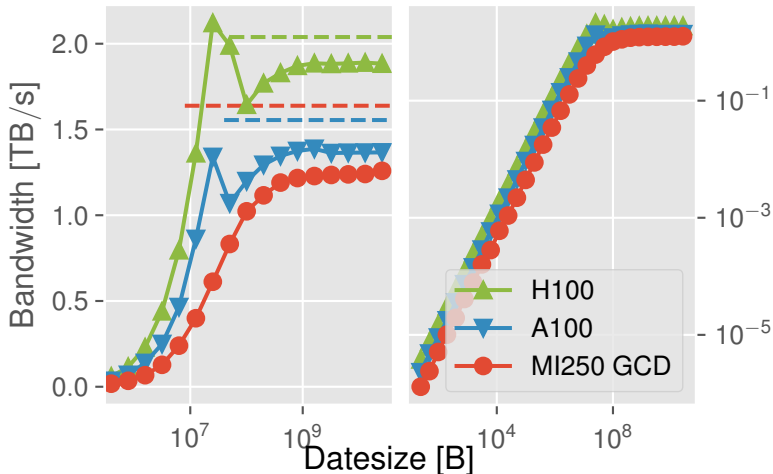
Staggered



CTE-ARM @ BSC, Kokkos 3.6, GCC 11.1, Clang 14.0

Nvidia A100, H100 and AMD MI250 (GPU)


Stream Triad



A100: JURECA, H100 & MI250: JURECA Evaluation Platform

Nvidia A100, H100 and AMD MI250 (GPU)

Staggered



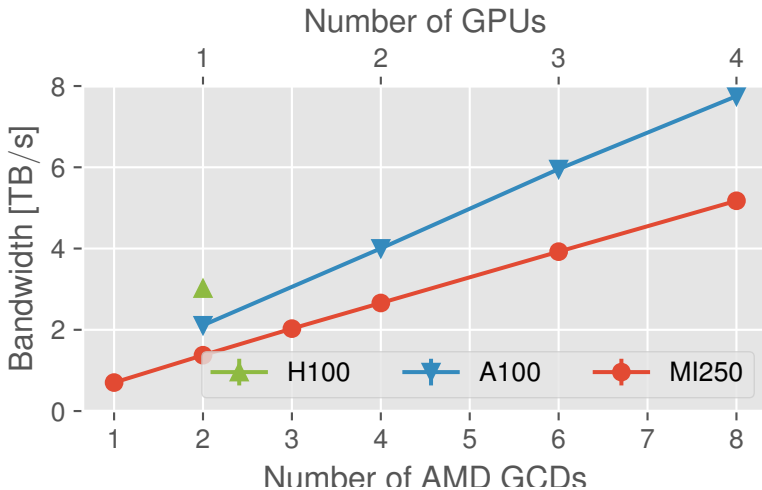
A100: JURECA @ JSC, Kokkos 4.0.1, GCC 11.3, CUDA 11.7, OpenMPI 4.1.4, LB 384,1

H100: JURECA Evaluation Platform @ JSC, Kokkos 4.0.1, GCC 11.3, CUDA 12.0, OpenMPI 4.1.4, LB 384,1

MI250: JURECA Evaluation Platform @ JSC, Kokkos 4.0.1, Clang 15, ROCm 5.4, OpenMPI 4.1.4

Nvidia A100, H100 and AMD MI250 (GPU)

Staggered



A100: JURECA, H100 & MI250: JURECA Evaluation Platform

END

Thank you for your attention!

